

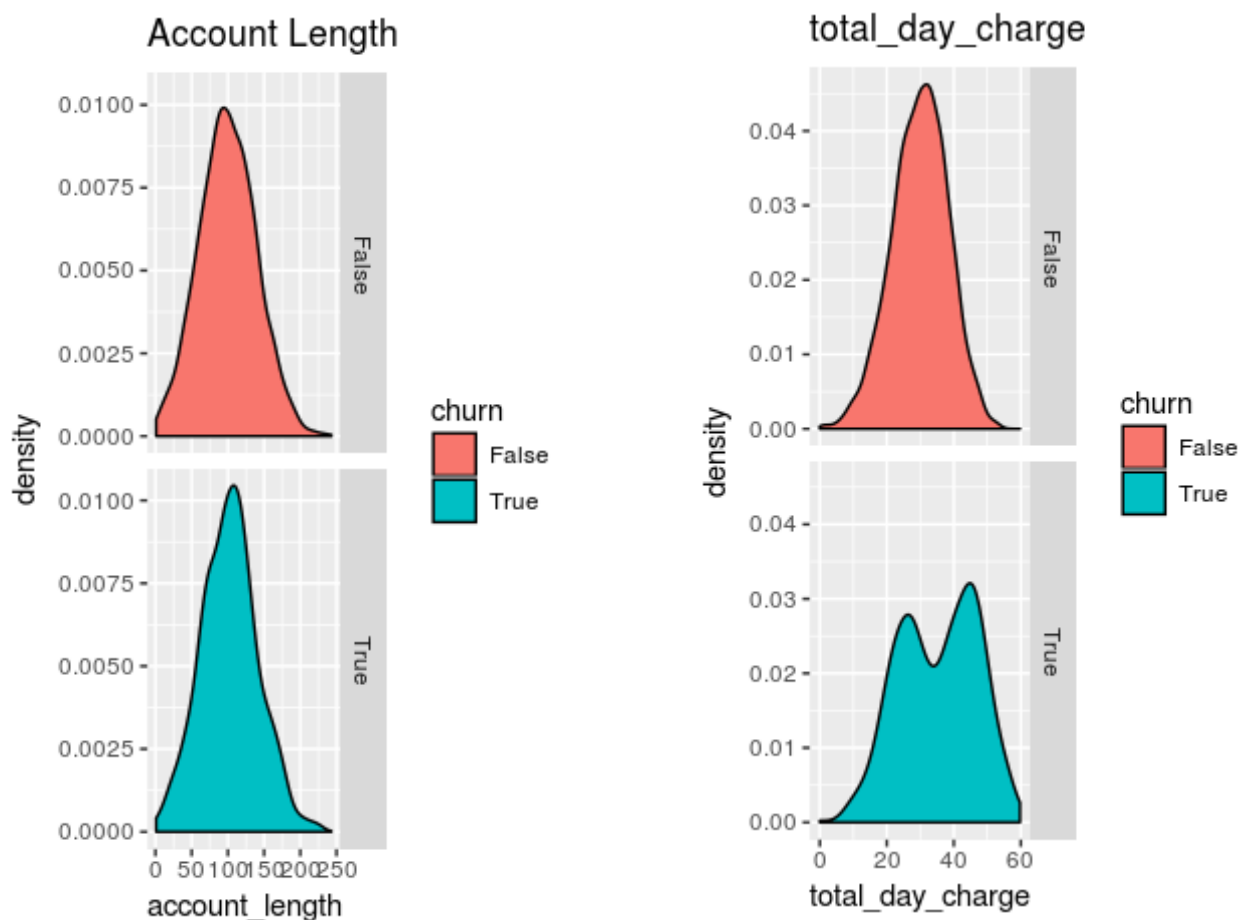
# Project Report: Assignment 3

## Customer Churn Prediction

Suyash Damle  
15CS10057

### Data Pruning/ Analysis of the attributes:

- The area code field is read into the data frame as numerical, where it should ideally be categorical. So, this field needs a change of type
- The data for phone numbers is more useful for the purpose of identification or verification of the provided data. Using it to train a classifier makes no sense, as it has *unique values for each observation*. This is verifiable by checking the number of level corresponding to the column “phone\_numbers”. This column has 5000 levels, for 5000 data values. Hence, the values are unique and useless for the classification task at hand. Thus, it is dropped out.
- Further checks are made during model training phase and required changes are made.



(Similar curves)

(Significantly different for some values)

(Some samples of the type of curves obtained: attribute value densities for both churn types)

The curves are roughly bell-shaped. In the case of account length, visually, the graphs are similar for both positive and negative churns. This implies that the **attribute account\_length may not have much**

**significance in the task of classification.** The same pattern is followed by a number of other attributes: total\_eve\_charge, total\_intl\_charge.

Some other attributes, such as the total\_day\_charge, have distinctly different curves in True churn values than the False values. So, they are more likely to have some significance at some point for the task of classification at hand.

## **Division of dataset**

The distribution of data in the **original data** is as follows:

| False | True |
|-------|------|
| 4293  | 707  |

The train to test ratio is kept at 8 : 2

After proper symmetrical division of the data, with equal proportions of True and False values in both test and train sets:

( Using the *createDataPartition()* function from the *caret* package)

**Train data:**

| False | True |
|-------|------|
| 3435  | 566  |

**Test data:**

| False | True |
|-------|------|
| 858   | 141  |

## **Evaluation metrics**

- The confusion metrics could be easily plotted and related information obtained, in each case, using the ***confusionMatrix()*** function.
- The parameters required for this are the test data classification and the predictions from the model.

***confusionMatrix ( predictions, dataTest\_y)***

- The function also returns the following important values:
  - Accuracy
  - Sensitivity – the Recall
  - Specificity (True negative rate) – this is what is important for us, as the True (churn) is being considered as the negative class
  - Pos. Predicted Value – the Precision

## **The Naive Bayes Classifier:**

Issues:

The naive bayes model takes *only categorical data*. So, to deal with the continuous numerical data abundant in this dataset, one needs to divide them into discrete classes. Two solutions exist:

1. Manual discretization : Deciding some fixed values for each attribute and grouping data above and below (and between) these values separately.

- Considering a probability density : One could assume that the attribute follows certain prob. distribution and then proceed:
  - assuming (say) gaussian distribution for all the attribute values in each class.
  - Eg: we may assume that the “total\_eve\_charge” in True and False class for all the observations follows normal distribution.
  - from the training data, finding the mean and standard deviation of the distribution for both, True and False classes
  - for each of the attribute to be considered henceforth, finding the probability density value at the observation data, corresponding to both True and False classes. The one with higher pdf value is considered to be the valid classification for the observation point.
  - So, if pdf is higher in False class, we put in category 1, else in category 2.

The used **naiveBayes()** function from **e1071** uses this second approach to deal with this kind of situation.

## Analysis from the Naive Bayes model

The naive bayes model, once trained could be used for obtaining further insight into the data attributes: the model begins with the calculations of posterior probabilities of the form:  $P(X|Y)$ , which are further expected to be useful in the calculation of  $P(Y|X)$  during the process of prediction.

The  $P(X|Y)$  data could be printed as:

### **model \$ tables:**

```
$state
state
dataTrain_y AK AL AR AZ CA CO CT DC DE FL GA HI IA
False 0.016351119 0.026104418 0.017785427 0.019219736 0.007458405 0.020080321 0.018646013 0.018646013 0.018072289 0.017785427 0.017211704 0.018932874 0.015777395
True 0.003241491 0.021069692 0.021069692 0.016207455 0.021069692 0.017828201 0.022690438 0.012965964 0.024311183 0.012965964 0.014586710 0.009724473 0.012965964

state
dataTrain_y ID IL IN KS KY LA MA MD ME MI MN MO MS
False 0.025817556 0.018359151 0.017498566 0.017785427 0.020940906 0.016637980 0.019219736 0.018932874 0.019219736 0.020940906 0.024956971 0.020654045 0.019506598
True 0.021069692 0.011345219 0.021069692 0.025931929 0.017828201 0.009724473 0.021069692 0.030794165 0.029173420 0.022690438 0.027552674 0.017828201 0.025931929

state
dataTrain_y MT NC ND NE NH NJ NM NV NY OH OK OR PA
False 0.018646013 0.017211704 0.020367183 0.016924842 0.019219736 0.020080321 0.018646013 0.017498566 0.022088353 0.023522662 0.016924842 0.021227768 0.015777395
True 0.027552674 0.021069692 0.014586710 0.011345219 0.016207455 0.034035656 0.016207455 0.022690438 0.027552674 0.024311183 0.017828201 0.027552674 0.011345219

state
dataTrain_y RI SC SD TN TX UT VA VT WA WI WV WY
False 0.019506598 0.019219736 0.018072289 0.017498566 0.020940906 0.023522662 0.025817556 0.021514630 0.017785427 0.022088353 0.028973035 0.024383247
True 0.011345219 0.022690438 0.014586710 0.017828201 0.037277147 0.024311183 0.008103728 0.014586710 0.027552674 0.014586710 0.030794165 0.011345219

$account_length
account_length
dataTrain_y [,1] [,2]
False 100.2786 39.64771
True 101.5689 39.85080
```

From this data, following conclusions are possible:

- The posterior probabilities,  $P(\text{state} | \text{churn})$  for most states do not differ much  
 Eg:  $P(\text{state} = \text{AK} | \text{churn} = \text{False}) \sim P(\text{state} = \text{AR} | \text{churn} = \text{False})$   
***This indicates that users of these states are all equally likely to remain with our services***
- In some states, the difference is significant  
 Eg:  $P(\text{state} = \text{AK} | \text{churn} = \text{True}) \gg P(\text{state} = \text{AL} | \text{churn} = \text{True})$   
***In such states, the user is more (or less) likely to continue with the /service as compared to other states***

```

$international_plan
      international_plan
dataTrain_y      no      yes
False 0.93628164 0.06371836
True   0.73239437 0.26760563

$voice_mail_plan
      voice_mail_plan
dataTrain_y      no      yes
False 0.7122491 0.2877509
True   0.8397887 0.1602113

$number_vmail_messages
      number_vmail_messages
dataTrain_y      [,1]      [,2]
False 8.351092 13.78803
True   5.021201 11.85650

$total_day_minutes
      total_day_minutes
dataTrain_y      [,1]      [,2]
False 175.6971 49.90298
True   206.5226 67.74762

$total_day_calls
      total_day_calls
dataTrain_y      [,1]      [,2]
False 99.86084 19.62373
True  100.75265 20.33670

$total_day_charge
      total_day_charge
dataTrain_y      [,1]      [,2]
False 29.86907 8.483467
True   35.10926 11.517274

$total_eve_minutes
      total_eve_minutes
dataTrain_y      [,1]      [,2]
False 198.9172 50.38457
True   212.8516 50.72032

```

and so on...

Final result:

#### Confusion Matrix and Statistics

|            | Reference |      |
|------------|-----------|------|
| Prediction | False     | True |
| False      | 829       | 75   |
| True       | 29        | 66   |

Accuracy : 0.8959

95% CI : (0.8753, 0.9141)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.0002898

Kappa : 0.5028

Mcnemar's Test P-Value : 1.021e-05

Sensitivity : 0.9662

Specificity : 0.4681

Pos Pred Value : 0.9170

Neg Pred Value : 0.6947

Prevalence : 0.8589

Detection Rate : 0.8298

Detection Prevalence : 0.9049

Balanced Accuracy : 0.7171

'Positive' Class : False

## The Decision Tree Classifier

- The decision tree classifier is trained using the **rpart()** (CART algorithm) function from the **e1071** package.
- Splitting is done on the basis of information gain and gini index – both

### Using information gain as a heuristic for splitting:

#### Confusion Matrix and Statistics

| Prediction | Reference |      |
|------------|-----------|------|
|            | False     | True |
| False      | 845       | 37   |
| True       | 13        | 104  |

Accuracy : 0.9499

95% CI : (0.9345, 0.9626)

No Information Rate : 0.8589

P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.7778

McNemar's Test P-Value : 0.001143

Sensitivity : 0.9848

Specificity : 0.7376

Pos Pred Value : 0.9580

Neg Pred Value : 0.8889

Prevalence : 0.8589

Detection Rate : 0.8458

Detection Prevalence : 0.8829

Balanced Accuracy : 0.8612

'Positive' Class : False

### Using Gini Index as heuristic:

#### Confusion Matrix and Statistics

| Prediction | Reference |      |
|------------|-----------|------|
|            | False     | True |
| False      | 852       | 45   |
| True       | 6         | 96   |

Accuracy : 0.9489

95% CI : (0.9334, 0.9618)

No Information Rate : 0.8589

P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.7619

McNemar's Test P-Value : 1.032e-07

Sensitivity : 0.9930

Specificity : 0.6809

Pos Pred Value : 0.9498

Neg Pred Value : 0.9412

Prevalence : 0.8589

Detection Rate : 0.8529

Detection Prevalence : 0.8979

Balanced Accuracy : 0.8369

'Positive' Class : False

So, *information gain is the choice, as it offers better specificity, which is what we need.*

- The created tree demonstrates which attributes give best result at the time of splitting. The visualization of the created tree has been attached with the codes and the report, in a pdf document.

## The Support Vector Machine (SVM) Models:

Two different kernels were tried :

### Linear

#### Confusion Matrix and Statistics

|            |     | Reference |      |
|------------|-----|-----------|------|
| Prediction |     | False     | True |
| False      | 858 | 129       |      |
| True       | 0   | 12        |      |

Accuracy : 0.8709

95% CI : (0.8485, 0.891)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.1477

Kappa : 0.1378

Mcnemar's Test P-Value : <2e-16

Sensitivity : 1.00000

Specificity : 0.08511

Pos Pred Value : 0.86930

Neg Pred Value : 1.00000

Prevalence : 0.85886

Detection Rate : 0.85886

Detection Prevalence : 0.98799

Balanced Accuracy : 0.54255

'Positive' Class : False

-> So, the linear SVM model (without scaling of attributes) *performs quite poorly in terms of the specificity.*

### Polynomial

- The polynomial kernel, having a large number of parameters offers the option of **tuning the parameters**. This technique lets the user enter the range of values (or, discrete values) to be checked.
- The system itself checks the classifier over all the values (*grid search*) and determine the optimal values.
- Two distinct runs of optimization were performed, for different parameter combinations :

#### Optimizing the polynomial degree:

```
model_svm_poly=tune.svm(churn~.,data=(dataTrain[2:21]),cost=5,kernel="polynomial",  
degree=c(2,3,4,5,6))
```

Optimal degree found : 2

## Confusion Matrix and Statistics

|            | Reference |      |
|------------|-----------|------|
| Prediction | False     | True |
| False      | 851       | 111  |
| True       | 7         | 30   |

Accuracy : 0.8819

95% CI : (0.8602, 0.9012)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.01866

Kappa : 0.2958

McNemar's Test P-Value : < 2e-16

Sensitivity : 0.9918

Specificity : 0.2128

Pos Pred Value : 0.8846

Neg Pred Value : 0.8108

Prevalence : 0.8589

Detection Rate : 0.8519

Detection Prevalence : 0.9630

Balanced Accuracy : 0.6023

'Positive' Class : False

## Optimizing degree of polynomial and the “gamma” parameter:

***model\_svm\_poly=tune.svm(churn~.,data=(dataTrain[2:21]),cost=5,kernel="polynomial",  
degree=c(2,3,4),gamma=c(0.01,0.1,1))***

Optimized parameters :      degree = 3  
                                 gamma = 0.1

## Confusion Matrix and Statistics

|            | Reference |      |
|------------|-----------|------|
| Prediction | False     | True |
| False      | 817       | 61   |
| True       | 41        | 80   |

Accuracy : 0.8979

95% CI : (0.8774, 0.916)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.0001368

Kappa : 0.5523

McNemar's Test P-Value : 0.0599338

Sensitivity : 0.9522

Specificity : 0.5674

Pos Pred Value : 0.9305

Neg Pred Value : 0.6612

Prevalence : 0.8589

Detection Rate : 0.8178

Detection Prevalence : 0.8789

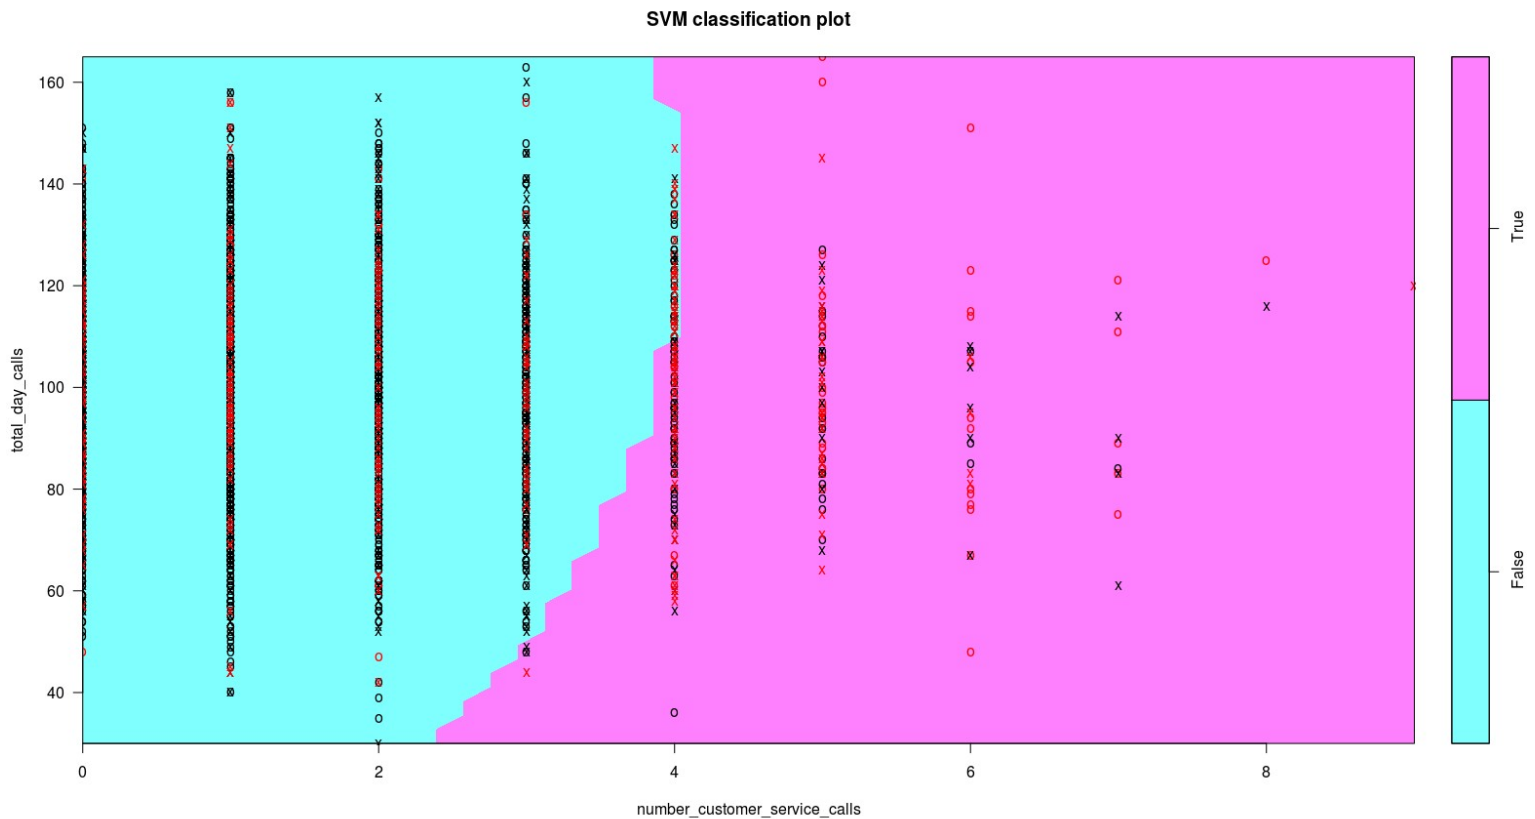
Balanced Accuracy : 0.7598

'Positive' Class : False

### Decision Boundary visualization:

After the training of the classifier, the decision boundary could be visualized wrt the training data using the function:

```
plot(model_svm_poly$best.model,data=dataTrain[2:21],formula=total_day_calls~number_customer_service_calls)
```



The variables in the x and y axis could be easily changed by simple replacement in the formula

- Running tuning process on more than 2 parameters takes a very long time to terminate and hence, SVM could not be checked and optimized any further. However, as the number of variations are large, some combination(s) may perform much better than the above results.



## Conclusion

The performances of all the explored models could be summarized as below:

| MODEL                                    | ACCURACY | RECALL  | PRECISION | SPECIFICITY |
|--|----------|---------|-----------|-------------|
| Naïve Bayes                              | 89.60%   | 96.60%  | 91.70%    | 46.80%      |
| Decision Tree (Binary, information gain) | 94.99%   | 98.48%  | 95.80%    | 73.76%      |
| SVM (Linear)                             | 87.90%   | 100.00% | 86.93%    | 8.50%       |
| SVM (Polynomial – degree = 2, gamma=1)   | 88.19%   | 99.18%  | 88.46%    | 21.28%      |
| SVM (Polynomial – degree = 3, gamma=0.1) | 89.79%   | 95.22%  | 93.05%    | 56.74%      |

Based on the above data, one could conclude that the **decision tree classifier with the information gain heuristic and the support vector machine with polynomial kernel perform the best** and the performance of SVM (presently slightly lagging behind) may beat it; however careful tuning of the parameters is required for that purpose.